

Web search engines

Paolo Ferragina

Dipartimento di Informatica

Università di Pisa

Two main difficulties

The Web:

Extracting “significant data” is difficult !!

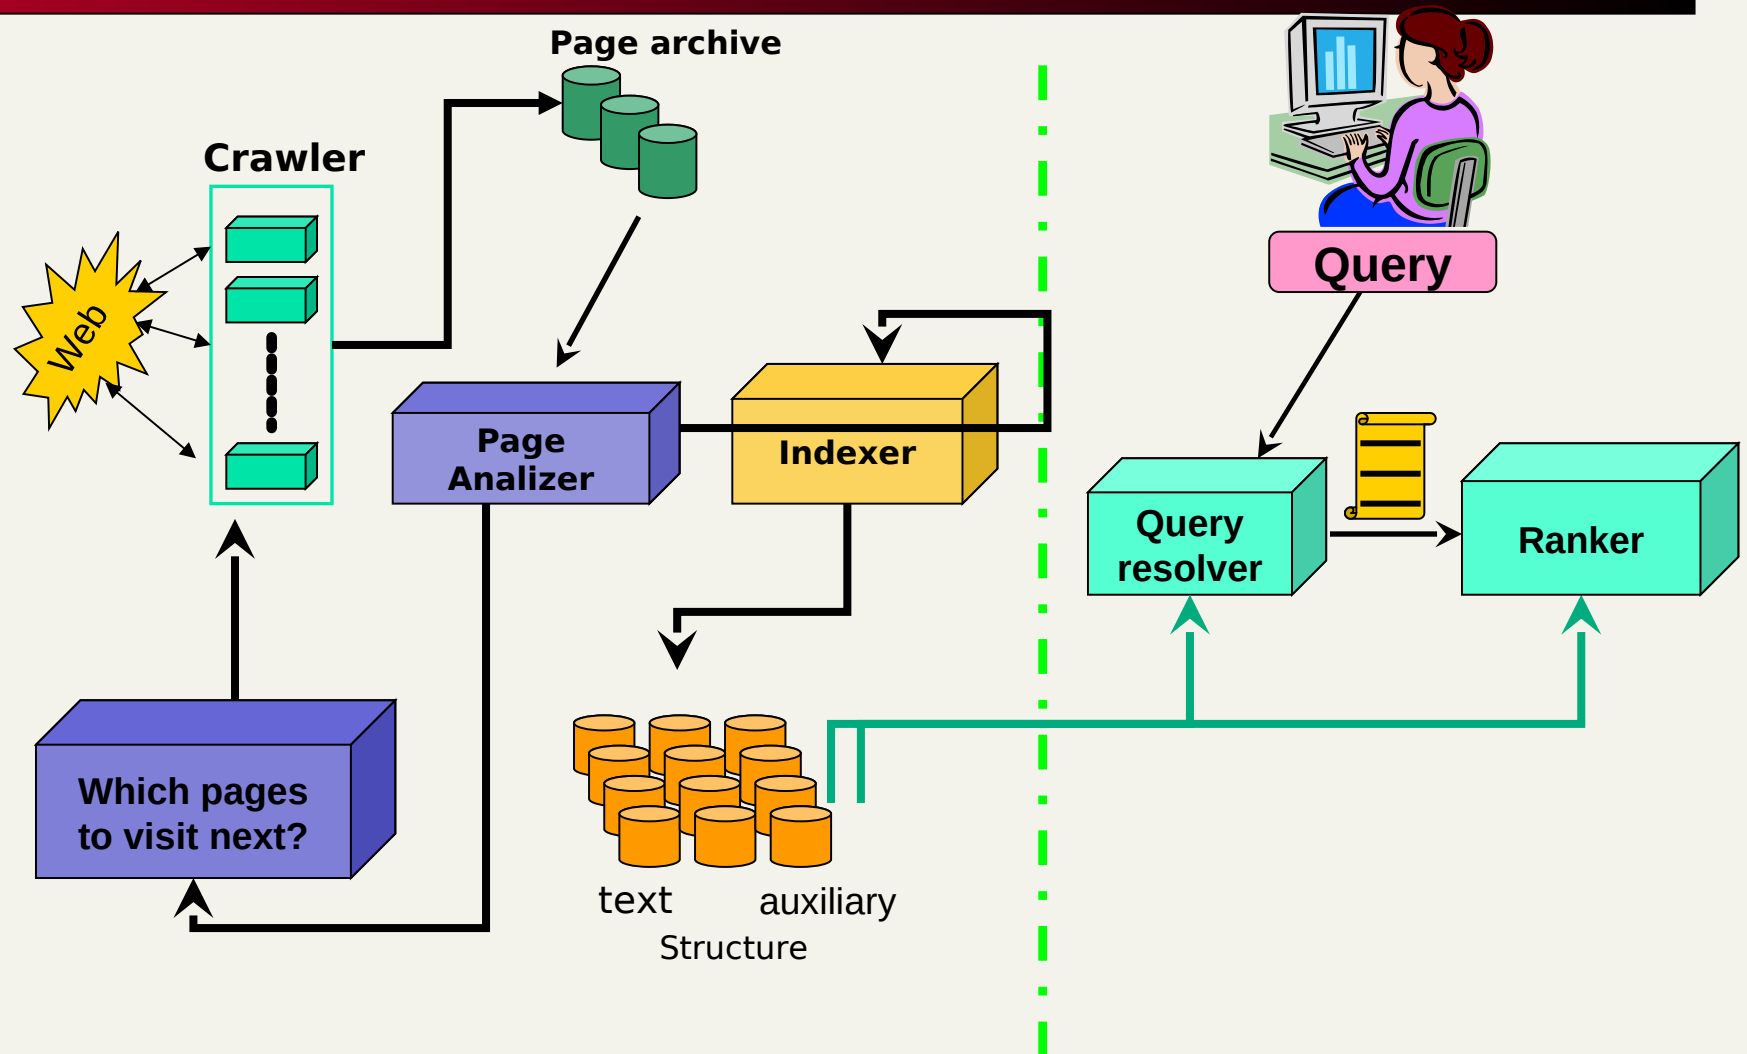
- **Size:** more than 1 trillion pages
- **Language and encodings:** hundreds...
- **Distributed authorship:** SPAM, format-less,...
- **Dynamic:** in one year 35% survive, 20% untouched

The User:

Matching “user needs” is difficult !!

- **Query composition:** short (2.5 terms avg) and imprecise
- **Query results:** 85% users look at just one result-page
- **Several needs:** Informational, Navigational, Transactional

The structure of a search Engine



The web graph: properties

Paolo Ferragina


Dipartimento di Informatica

Università di Pisa

The Web's Characteristics

- It's a graph whose size is
 - 1 trillion of pages is available
 - 50 billion pages crawled (09/15)
 - 5-40K per page => terabytes & terabytes
 - Size grows every day!!
 - Actually the web is infinite... Calendars...
- It's a dynamic graph with
 - 8% new pages, 25% new links change weekly
 - Life time of about 10 days

The Bow Tie



Crawling


Paolo Ferragina

Dipartimento di Informatica

Università di Pisa

Spidering


- 24h, 7days “walking” over the web graph
- Recall that:
 - Direct graph $G = (N, E)$
 - N changes (insert, delete) \gg trillion nodes
 - E changes (insert, delete) > 10 links per node
 - Trillion entries in posting lists
 - Many more if we consider also the word positions in every document where it occurs.




Crawling Issues

- How to crawl?
 - *Quality*: “Best” pages first
 - *Efficiency*: Avoid duplication (or near duplication)
 - *Etiquette*: Robots.txt, Server load concerns (Minimize load)
 - *Malicious pages*: Spam pages, Spider traps – incl dynamically generated
- How much to crawl, and thus index?
 - *Coverage*: How big is the Web? How much do we cover?
 - *Relative Coverage*: How much do competitors have?
- How often to crawl?
 - *Freshness*: How much has changed?
 - *Frequency*: Commonly insert time gap btw host requests

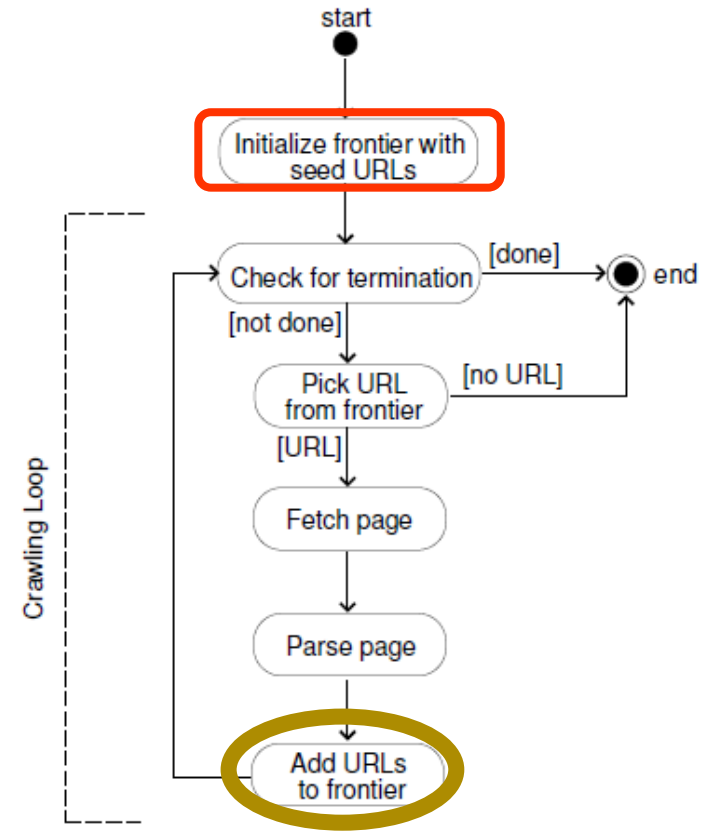

Crawling picture




Updated crawling picture



A small example



Crawler “life cycle”



One Link Extractor per page:

```
while(<Page Repository is not empty>){  
  <take a page p (check if it is new)>  
  <extract links contained in p within href>  
  <extract links contained in javascript>  
  <extract .....>  
  <insert these links into the Priority Queue>  
}
```

One Downloader per page:

```
while(<Assigned Repository is not empty>){  
  <extract url u>  
  <download page(u)>  
  <send page(u) to the Page Repository>  
  <store page(u) in a proper archive,  
    possibly compressed>  
}
```

One single Crawler Manager:

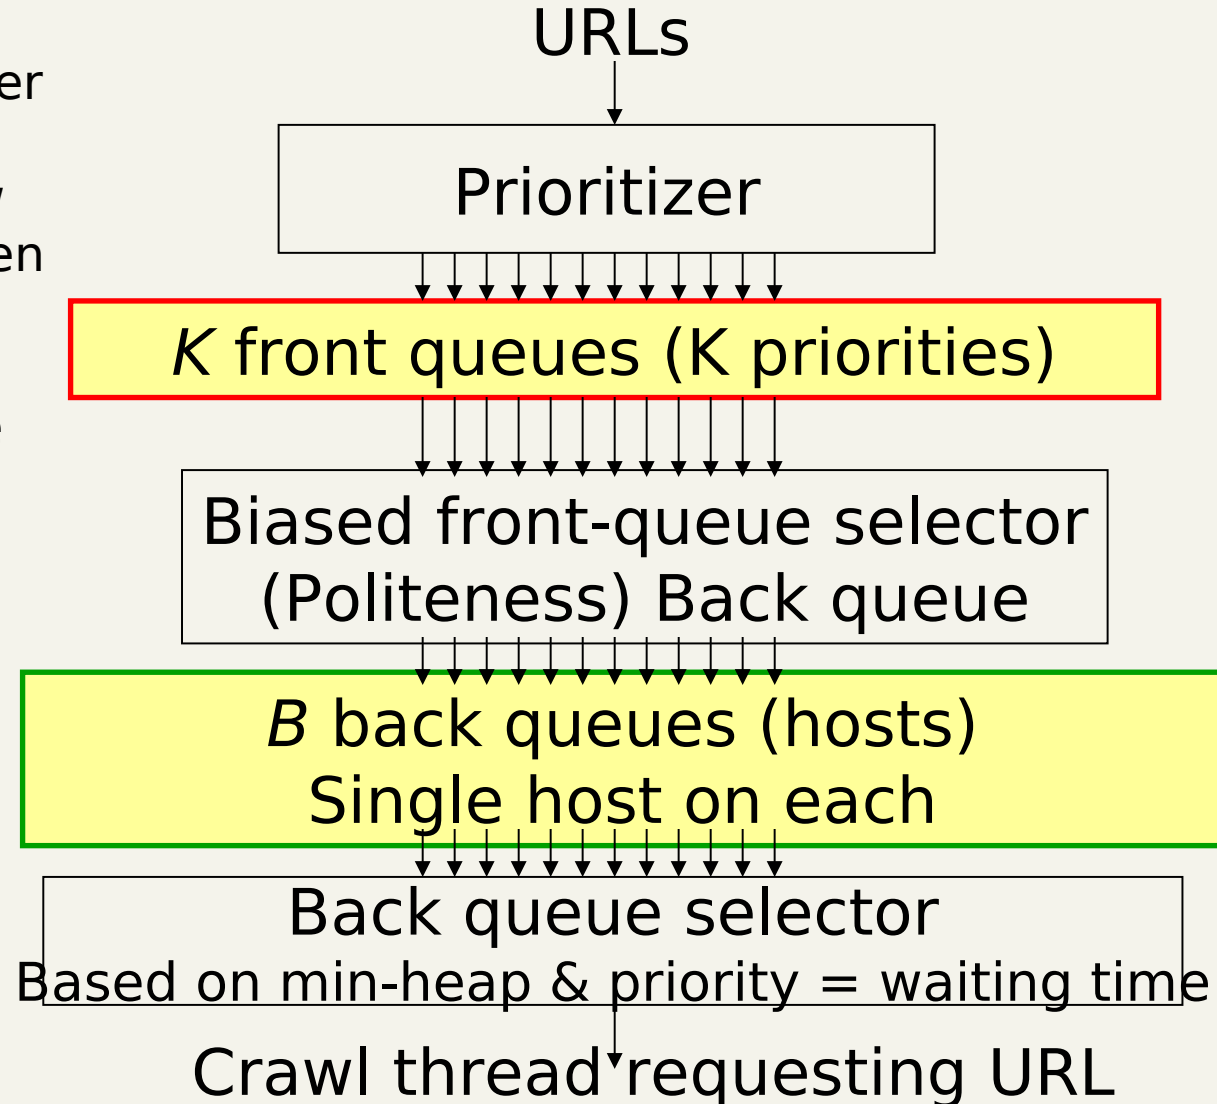
```
while(<Priority Queue is not empty>){  
  <extract some URL u having the highest priority>  
  foreach u extracted {  
    if ( (u ∉ “Already Seen Page” ) ||  
        (u ∈ “Already Seen Page” && <u’s version on the Web is more recent> )  
      ) {  
      <resolve u wrt DNS>  
      <send u to the Assigned Repository>  
    }  
  }  
}
```

URL frontier visiting

- Given a page P, define how “good” P is.
- Several *metrics* (via priority assignment):
 - BFS, DFS, Random
 - Popularity driven (*PageRank*, full vs partial)
 - Topic driven or focused crawling
 - Combined
- How to fast check whether the URL is new ?

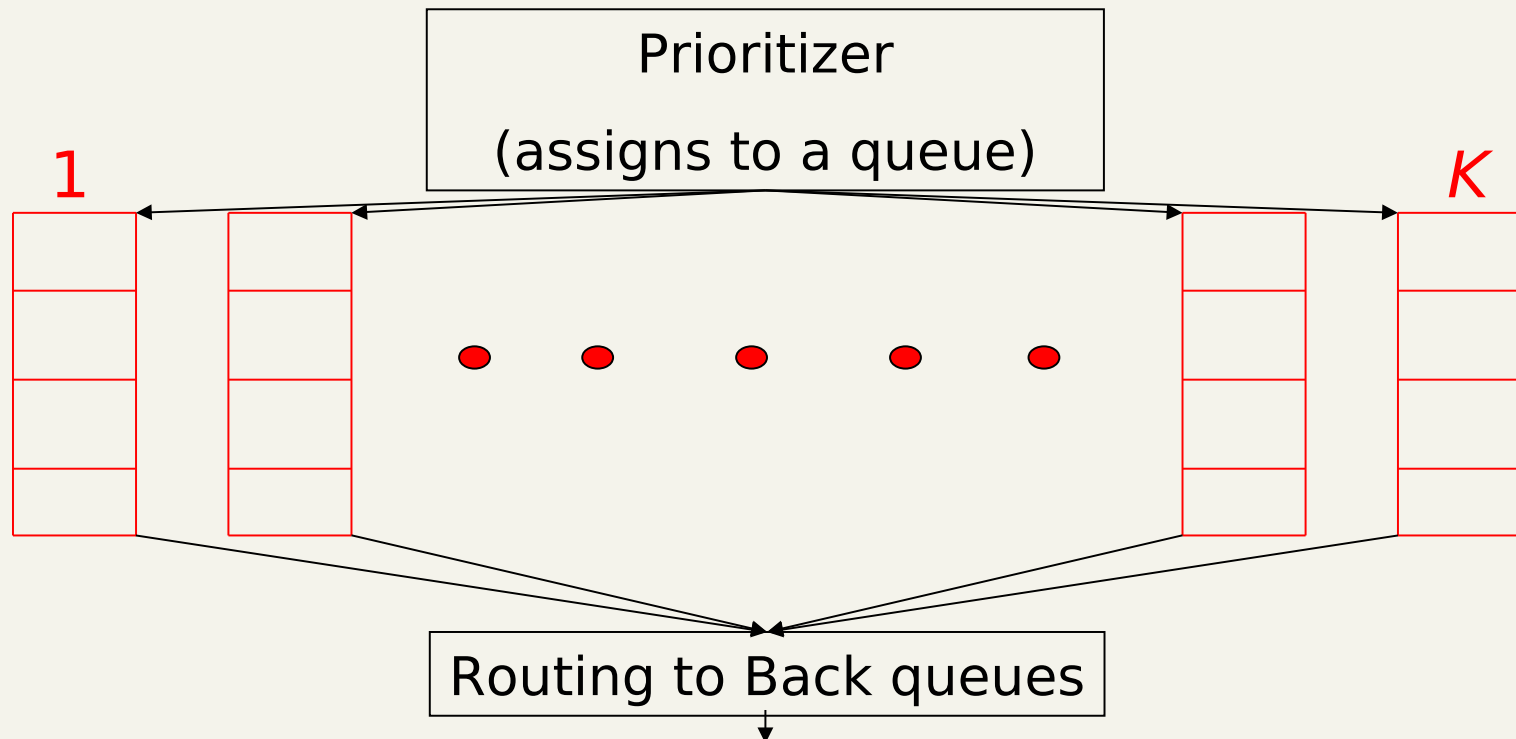
Mercator

1. Only one connection per host is open at a time;
2. a waiting time of a few seconds occurs between successive requests to the same host;
3. high-priority pages are crawled preferentially.




Front queues


- **Front queues** manage **prioritization**:
 - Prioritizer assigns to an URL an integer priority (refresh, quality, application specific) between 1 and K
 - Appends URL to corresponding queue, according to priority



Back queues




Back queue request → Select a **front** queue *randomly*, biasing towards higher queues



Back queues

- **Back queues** enforce **politeness**:
 - Each back queue is kept non-empty
 - Each back queue contains only URLs from a **single** host

Back queue request → Select a **front** queue *randomly*, biasing towards higher queues



The **min-heap**

- It contains one entry per back queue
- The entry is **the earliest time t_e** at which the host corresponding to the back queue can be “hit again”
- This earliest time is determined from
 - Last access to that host
 - Any time buffer heuristic we choose

The crawl thread

- A crawler seeks a URL to crawl:
 - **Extracts the root of the heap:** So it is an URL at the head of some back queue q (*and then removes it*)
 - **Waits** the indicated time t_{url}
 - **Parses URL** and adds its out-links to the Front queues
- If back queue q **gets empty**, pulls a **URL v** from some front queue (*more prob for higher queues*)
 - If there's already a back queue for v 's host, append v to it and repeat until q gets not empty;
 - Else, make q the back queue for v 's host
- If back queue q **is non-empty**, pick URL and add it to the min-heap with priority = waiting time t_{url}

Keep crawl threads busy ($B = 3 \times \text{threads}$)