

# Locality-sensitive hashing and its applications

Paolo Ferragina  
University of Pisa

ACM Kanellakis Award  
2013

# A frequent issue

Given  $U$  users, described with a set of  $d$  features, the goal is to find (the largest) **group** of *similar users*

**Features** = Personal data, preferences, purchases, navigational behavior, search behavior, followers/ing, ...

**Similarity( $u_1, u_2$ )** is a function that, taken the set of features of users  $u_1$  and  $u_2$ , returns a value in  $[0,1]$



# Solution #1

Try all *groups of users* and, for each group, check the (average) similarity among all its users.

# Sim computations  $\cong 2^U \times U^2$

In the case of Facebook this is  $> 2^{1\text{billion}} \times (10^9)^2$

If we limit groups to have a size  $\leq L$  users

# Sim computations  $\cong U^L \times L^2$

(Even if 1ns/sim and  $L=10$ , it is  $> (10^9)^{10} / 10^9$  secs  $> 10^{70}$  years)

No faster CPU/GPU, multi-cores,... could help !

# Solution #2: introduce approximation

Interpret every user as a point in a **d**-dim space, and then apply a **clustering** algorithm

**K-means**

Pick  $K=2$  centroids at random

Compute clusters

Re-determine centroids

Re-compute clusters

Re-determine centroids

Re-compute clusters

Converged!

Each iteration takes  $\cong K \times U$

# Solution #2: few considerations

- Cost per iteration =  $\mathbf{K} \times \mathbf{U}$ , #iterations is typically small
- What about optimality ?** It is locally optimal [recently, some researchers showed how to introduce some guarantee]
- What about the Sim-cost ?** Comparing users/points costs  $\Theta(d)$  in time and space [notice that  $d$  may be

bi/mi

In  $T$  time, we can manage  $U = T^{1/3}$

- What users**

$I_3 < U_L$

[ $\cong y$ ]

Using  $s$ -faster CPU  $\approx$  using  $sT$  time an old CPU

$\rightarrow$  we can manage  $(c * T)^{1/3} = c^{1/3} T^{1/3}$  users

## Solution #3: introduce randomization

Generate a **fingerprint** for every user that is **much shorter** than **d** and allows to transform *similarity* into ***equality*** of fingerprints.

- ✓ It is ***randomized***, correct ***with high probability***
- ✓ It guarantees ***local access*** to data, which is good for speed in disk/distributed setting

ACM Kanellakis Award  
2013

# A warm-up problem

---

- Consider vectors  $p, q$  of  $d$  binary features
- Hamming distance

$D(p, q) = \text{\#bits where } p \text{ and } q \text{ differ}$

- Define hash function  $h$  by choosing a set  $l$  of  $k$  random coordinates

$h(p) = \text{projection of vector } p \text{ on } l\text{'s coordinates}$

Example: Pick  $l = \{1, 4\}$  ( $k=2$ ), then  $h(p = \mathbf{0}10\mathbf{1}1) = 01$

# A key property

$\Pr[\text{picking } x \text{ s.t. } p[x]=q[x]] = (d - D(p,q))/d$

$$\Pr[h(p) = h(q)] = \left(1 - \frac{D(p,q)}{d}\right)^k$$

We can vary the probability by changing



*p versus q*



# ■ =  $D(p,q)$

# ■ =  $d - D(p,q)$

$= s^k$   
where  $s$  is the  
similarity  
between  $p$  and  $q$



Larger  $k$   
Smaller False Positive  
What about False Negatives?



Larger L  
Smaller False Negatives

# Reiterate L times

- 1) Repeat **L times** the **k-projections**  $h_i(p)$
- 2) We set  $g(p) = \langle h_1(p), h_2(p), \dots, h_L(p) \rangle$  **Sketch(p)**
- 3) Declare «p matches q» if **at least** one  $h_i(p)=h_i(q)$

## Example:

We set  $k=2$ ,  $L=3$ , let  $p = 01\mathbf{0}01$  and  $q = 01\mathbf{1}01$

- $I_1 = \{3,4\}$ , we have  $h_1(p) = 00$  and  $h_1(q)=10$
- $I_2 = \{1,3\}$ , we have  $h_2(p) = 00$  and  $h_2(q)=01$
- $I_3 = \{1,5\}$ , we have  $h_3(p) = 01$  and  $h_3(q)=01$

p and q declared  
to match !!


# Measuring the error probability

$$\Pr[h_i(p) = h_i(q)] = \left(1 - \frac{D(p, q)}{d}\right)^k = s^k$$

The  $g()$  consists of  $L$  independent hashes  $h_i$

$$\Pr[g(p) \text{ matches } g(q)] = 1 - \Pr[h_i(p) \neq h_i(q), \forall i=1, \dots,$$


$$L] \quad \Pr(g(p) \cong g(q)) = 1 - (1 - s^k)^L$$



# The case: Groups of similar items

Buckets provide the candidate similar items  
«Merge» similar sets over  $L$  rounds if they share items

**No Tables !  
→ SORT**




If  $p \approx q$ , then they fall in at least one same bucket

# The case of *on-line query*

Given a query  $w$ , find the similar indexed vectors:

check the vectors in the buckets  $h_j(w)$  for all  $j=1, \dots, L$



# LSH *versus* K-means

- **What about optimality ?** K-means is locally optimal [LSH finds correct clusters with high probability]
- **What about the Sim-cost ?** K-means compares vectors of  $d$  components [LSH compares very short (sketch) vectors]
- **What about the cost per iteration?** Typically K-means requires few iterations, each costs  $K \times U \times d$  [LSH sorts  $U$  short items, few scans]
- **What about  $K$  ?** In principle  $K$  can be arbitrary, ...,  $U$  [LSH does not need to know  $K$ ]

You could apply K-means over LSH-sketch vectors !!