

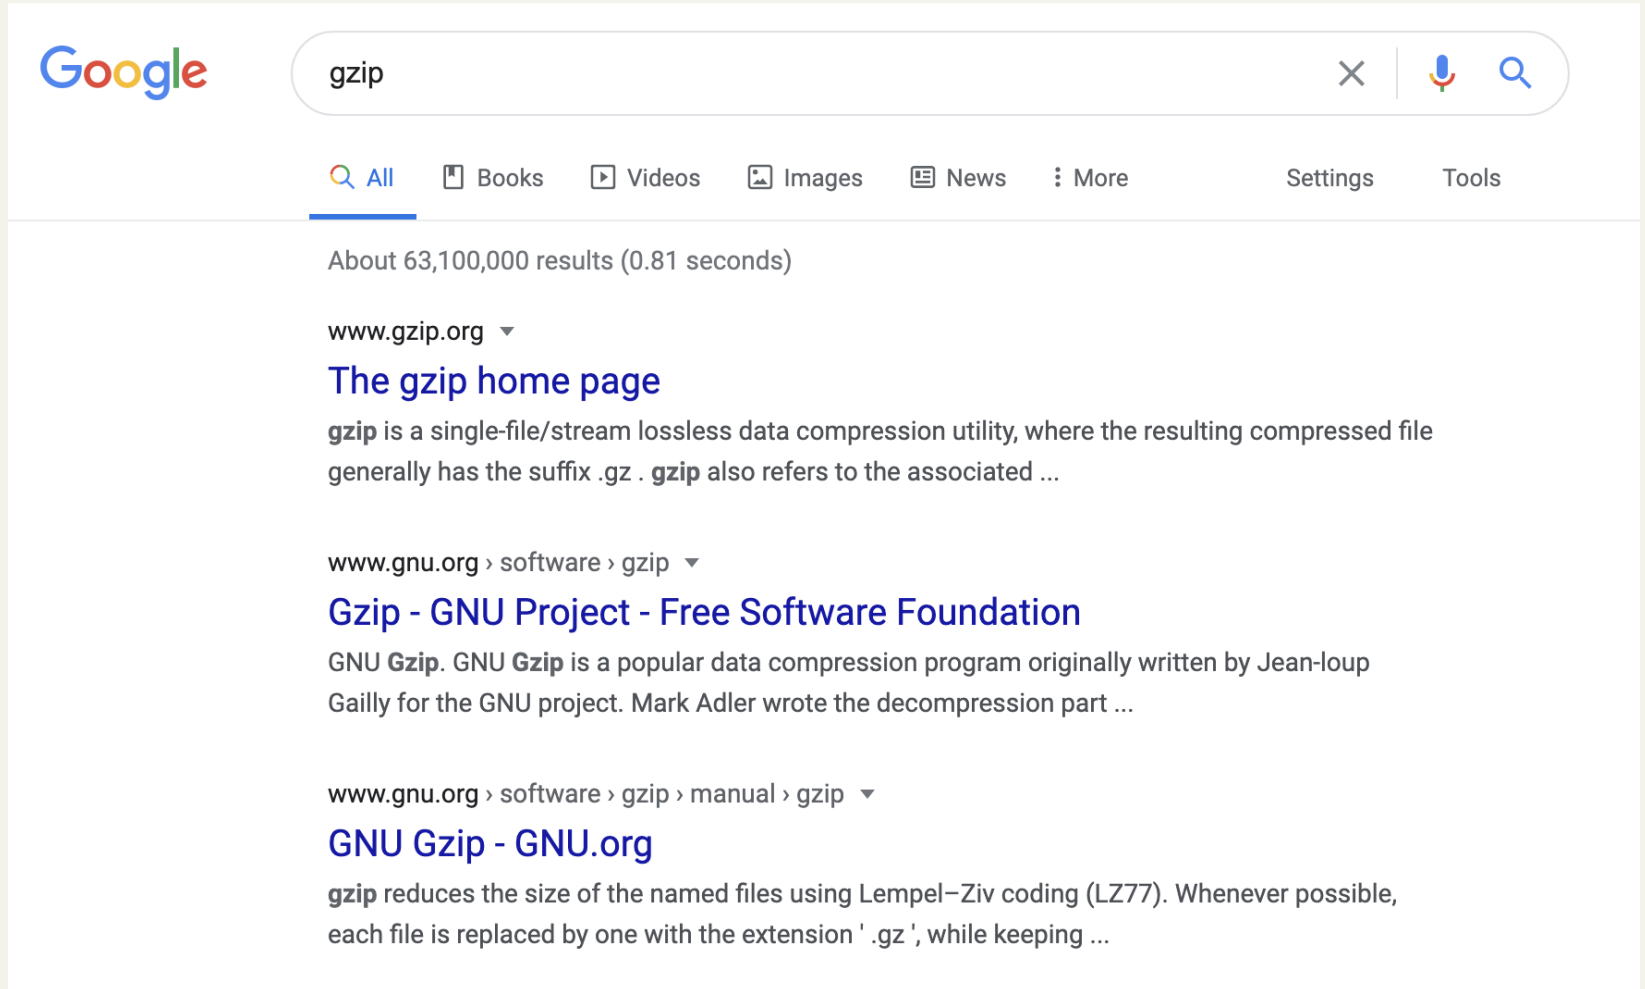
# Compression of documents

Paolo Ferragina

Dipartimento di Informatica

Università di Pisa



# Raw docs are needed









A screenshot of a Google search interface. The search bar at the top contains the text "gzip". To the left of the search bar is the Google logo. To the right are icons for voice search and image search. Below the search bar, there are tabs for "All", "Books", "Videos", "Images", "News", and "More". The "All" tab is selected. Below the tabs, the search results are displayed. The first result is from "www.gzip.org" and is titled "The gzip home page". The description says "gzip is a single-file/stream lossless data compression utility, where the resulting compressed file generally has the suffix .gz. gzip also refers to the associated ...". The second result is from "www.gnu.org" and is titled "Gzip - GNU Project - Free Software Foundation". The description says "GNU Gzip. GNU Gzip is a popular data compression program originally written by Jean-loup Gailly for the GNU project. Mark Adler wrote the decompression part ...". The third result is also from "www.gnu.org" and is titled "GNU Gzip - GNU.org". The description says "gzip reduces the size of the named files using Lempel–Ziv coding (LZ77). Whenever possible, each file is replaced by one with the extension '.gz', while keeping ...".

Google

gzip

× |  

 All  Books  Videos  Images  News  More Settings Tools

About 63,100,000 results (0.81 seconds)

www.gzip.org ▼

**The gzip home page**

gzip is a single-file/stream lossless data compression utility, where the resulting compressed file generally has the suffix .gz . gzip also refers to the associated ...

www.gnu.org › software › gzip ▼


**Gzip - GNU Project - Free Software Foundation**

GNU Gzip. GNU Gzip is a popular data compression program originally written by Jean-loup Gailly for the GNU project. Mark Adler wrote the decompression part ...

www.gnu.org › software › gzip › manual › gzip ▼

**GNU Gzip - GNU.org**

gzip reduces the size of the named files using Lempel–Ziv coding (LZ77). Whenever possible, each file is replaced by one with the extension '.gz', while keeping ...



A fast compressor/decompressor

 Search projects

[Project Home](#)
[Downloads](#)
[Wiki](#)
[Issues](#)
[Source](#)
[Summary](#)
[People](#)

## Project Information

 +76 Recommend this on Google

[Project feeds](#)
**Code license**
[New BSD License](#)
**Labels**
[Google](#), [Compression](#)

**Members**
[se...@google.com](#)
[1 committer](#)

## Featured


**Downloads**
[snappy-1.0.5.tar.gz](#)
[Show all »](#)

## Links

**Groups**
[General Snappy discussion](#)

Snappy is a compression/decompression library. It does not aim for maximum compression, or compatibility with any other compression library; instead, it aims for very high speeds and reasonable compression. For instance, compared to the fastest mode of zlib, Snappy is an order of magnitude faster for most inputs, but the resulting compressed files are anywhere from 20% to 100% bigger. On a single core of a Core i7 processor in 64-bit mode, Snappy compresses at about 250 MB/sec or more and decompresses at about 500 MB/sec or more.

Snappy is widely used inside Google, in everything from BigTable and MapReduce to our internal RPC systems. (Snappy has previously been referred to as "Zippy" in some presentations and the likes.)

For more information, please see the [README](#). Benchmarks against a few other compression libraries (zlib, LZO, LZF, FastLZ, and QuickLZ) are included in the source code distribution. The source code also contains a [formal format specification](#), as well as a [specification for a framing format](#) useful for higher-level framing and encapsulation of Snappy data, e.g. for transporting Snappy-compressed data across HTTP in a streaming fashion. Note that there is currently no known code implementing the latter.

Snappy is written in C++, but C bindings are included, and several bindings to other languages are maintained by third parties:

- [C89 port](#)
- [Common Lisp](#)
- Erlang: [esnappy](#), [snappy-erlang-nif](#)
- [Go](#)
- [Haskell](#)
- Java: [JNI wrapper](#), [native reimplementation](#)
- [Node.js](#)
- [Perl](#)
- [PHP](#)
- [Python](#)
- [Ruby](#)

If you know of more, do not hesitate to let us know. The easiest way to get in touch is via the [Snappy discussion mailing list](#).

## Introducing Brotli: a new compression algorithm for the internet

Posted: Tuesday, September 22, 2015



At Google, we think that internet users' time is valuable, and that they shouldn't have to wait long for a web page to load. Because fast is better than slow, two years ago we published the [Zopfli compression algorithm](#). This received such positive feedback in the industry that it has been integrated into many compression solutions, ranging from PNG optimizers to preprocessing web content. Based on its use and other modern compression needs, such as [web font compression](#), today we are excited to announce that we have developed and open sourced a new algorithm, the [Brotli compression algorithm](#).


While Zopfli is [Deflate](#)-compatible, Brotli is a whole new [data format](#). This new format allows us to get 20–26% higher compression ratios over Zopfli. In our study '[Comparison of Brotli, Deflate, Zopfli, LZMA, LZHAM and Bzip2 Compression Algorithms](#)' we show that Brotli is roughly as fast as [zlib's](#) Deflate implementation. At the same time, it compresses slightly more densely than [LZMA](#) and [bzip2](#) on the [Canterbury corpus](#). The higher data density is achieved by a 2nd order context modeling, re-use of entropy codes, larger memory window of past data and joint distribution codes. Just like Zopfli, the new algorithm is named after Swiss bakery products. Brötli means 'small bread' in Swiss German.

The smaller compressed size allows for better space utilization and faster page loads. We hope that this format will be supported by major browsers in the near future, as the smaller compressed size would give additional benefits to mobile users, such as lower data transfer fees and reduced battery use.

By Zoltan Szabadka, Software Engineer, Compression Team



 Labels

 Archive

 Feed

# Apple Open-Sources its New Compression Algorithm LZFSE

by [Sergio De Simone](#) on Jul 02, 2016 | [Discuss](#)

Share [+](#) [Twitter](#) [Y](#) [Reddit](#) [Facebook](#) [Email](#)

[My Reading List](#)

[Read later](#)

Apple has open-sourced its new lossless compression algorithm, [LZFSE](#), introduced last year with iOS 9 and OS X 10.10. According to Apple, LZFSE provides the same compression gain as ZLib level 5 while being 2x–3x faster and with higher energy efficiency.

LZFSE is based on Lempel-Ziv and uses [Finite State Entropy coding](#), based on Jarek Duda's work on [Asymmetric Numeral Systems](#) (ANS) for entropy coding. Shortly, ANS aims to “end the trade-off between speed and rate” and can be used both for precise coding and very fast encoding, with support for data encryption. LZFSE is one of a [growing number](#) of compression libraries that use ANS in place of the more traditional [Huffman](#) and [arithmetic coding](#).

Admittedly, LZFSE does not aim to be the best or fastest algorithm out there. In fact, Apple states that [LZ4](#) is faster than LZFSE while [LZMA](#) provides a higher compression ratio, albeit at the cost of being an order of magnitude slower than other options available in Apple SDKs. LZFSE is Apple's suggested option when compression and speed are more or less equally important and you want reduce energy consumption.

LZFSE reference implementation is available on [GitHub](#). Building on macOS is as easy as executing:

```
$ xcodebuild install DSTROOT=/tmp/lzfse.dst
```

If you want to build LZFSE for a current iOS device, you can execute:

```
xcodebuild -configuration "Release" -arch armv7 install DSTROOT=/tmp/lzfse.dst
```

## RELATED CONTENT

[Facebook Open-Sources New Compression Algorithm Outperforming Zlib](#) Sep 02, 2016

[Swift 3 is Out](#) Sep 20, 2016

[Image Processing for iOS](#) Jul 05, 2016



[Oracle Gives NetBeans to the Apache Foundation](#) Sep 19, 2016



[Open Source Swift Under the Hood](#) Apr 09, 2016



[IBM's Swift on the Server](#) Apr 09, 2016



# LZ77



Algorithm's step:

- Output  $\langle \text{dist}, \text{len}, \text{next-char} \rangle$
- Advance by  $\text{len} + 1$

A buffer "window" has fixed length and moves

$\langle 0, 0, a \rangle \langle 1, 1, c \rangle \langle 3, 4, b \rangle \dots$

# Example: LZ77 with window

| a a c a a c a b c a b a a a c (0, 0, a)

a | a c a a c a b c a b a a a c (1, 1, c)

a a c | a a c a b c a b a a a c (3, 4, b)

a a c a a c a b | c a b a a a c (3, 3, a)

a a c a a c | a b c a b a a a c (1, 2, c)

Window size = 6

Longest match

Next character

*within W*

Gzip -1...-9

Which is faster in decompression  
among gzip -1...-9 ?

# LZ77 Decoding

Decoder keeps same dictionary window as encoder.

- Finds substring  $\langle \text{len}, \text{dist}, \text{char} \rangle$  in previously decoded text
- Inserts a copy of it

What if  $\text{len} > \text{dist}$  ? (overlap with text to be compressed)

- E.g. seen = abcd, next codeword is  $\langle 2, 9, e \rangle$

```
for (i = 0; i < len; i++)  
    out[cursor+i] = out[cursor-d+i]
```

- Output is correct: **abcd**cdcdcdcdce



You find this at: [www.gzip.org/zlib/](http://www.gzip.org/zlib/)



# Squash Compression Benchmark

The Squash library is an abstraction layer for compression algorithms, making it trivial to switch between them... or write a benchmark which tries them all, which is what you see here!

The Squash Compression Benchmark currently consists of 28 datasets, each of which is tested against 29 plugins containing 46 codecs at every compression level they offer—the number varies by codec, but there are 235 in total, yielding 6,580 different settings. The benchmark is currently run on 9 different machines for a current grand total of 59,220 configurations, and growing.

[SKIP TO RESULTS \(PRETTY PICTURES!\)](#)[LEARN MORE ABOUT SQUASH ↗](#)

# Configuration

## Choose a dataset

Different codecs can behave *very* differently with different data. Some are great at compressing text but horrible with binary data, some excel with more repetitive data like logs. Many have long initialization times but are fast once they get started, while others can compress/decompress small buffers almost instantly.

This benchmark is run against many standard datasets. Hopefully one of them is interesting for you, but if not don't worry—you can use Squash to easily run your own benchmark with whatever data you want. That said, if you think you have a somewhat common use case, please [let us know](#)—we may be interested in adding the data to this benchmark.

### Note

The default dataset is selected randomly.

Name ^	Source	Description	Size
<input type="radio"/> <b>alice29.txt</b>	<a href="#">Canterbury Corpus</a>	English text	148.52 KiB
<input type="radio"/> <b>asyoulik.txt</b>	<a href="#">Canterbury Corpus</a>	Shakespeare	122.25 KiB
<input type="radio"/> <b>cp.html</b>	<a href="#">Canterbury Corpus</a>	HTML source	24.03 KiB
<input type="radio"/> <b>dickens</b>	<a href="#">Silesia Corpus</a>	Collected works of Charles Dickens	9.72 MiB
<input type="radio"/> <b>enwik8</b>	<a href="#">Large Text Compression Benchmark</a>	The first 10 <sup>8</sup> bytes of the English Wikipedia dump on Mar. 3, 2006	95.37 MiB
<input type="radio"/> <b>fields.c</b>	<a href="#">Canterbury Corpus</a>	C source	10.89 KiB
<input type="radio"/> <b>fireworks.jpeg</b>	<a href="#">Snappy</a>	A JPEG image	120.21 KiB
<input type="radio"/> <b>geo.protodata</b>	<a href="#">Snappy</a>	A set of Protocol Buffer data	115.81 KiB










# Choose a machine

If you think certain algorithms are always faster, you've got another thing coming! Different CPUs can behave very differently with the same data.


The Squash benchmark is currently run on many of the machines I have access to—this happens to be fairly recent Intel CPUs, and a mix of ARM SBCs. There is [an entry in the FAQ](#) with more details.

## Note

The default machine is selected randomly.

Name ^	Status	CPU/SoC	Architecture	Clock Speed	Memory	Platform	Distro	Kernel	Compiler	CSV
<input type="radio"/> beagleboard-xm	✓	Texas Instruments DM3730	armv7l	1 GHz	512 MiB	BeagleBoard-xM revision B	Ubuntu 15.04	4.1.5	gcc-4.9.2	
<input type="radio"/> e-desktop	✓	Intel® Core™ i3-2105	x86_64	3.1 GHz	8 GiB	Asus P8H61-H	Fedora 22	4.1.4	gcc-5.1.1	
<input type="radio"/> hoplite	✓	Intel® Core™ i7-2630QM	x86_64	2 GHz	6 GiB	Toshiba Satellite A660-X	Fedora 22	4.1.4	gcc-5.1.1	
<input type="radio"/> odroid-c1	✓	Amlogic S805	armv7l	1.5 GHz	1 GiB	ODROID-C1	Ubuntu 14.04.4	3.10.80	gcc-4.9.2	
<input checked="" type="radio"/> peltast	✓	Intel® Xeon® Processor E3-1225 v3	x86_64	3.2 GHz	20 GiB	Lenovo ThinkServer TS140	Fedora 22	4.1.6	gcc-5.1.1	
<input type="radio"/> phalanx	✓	Intel® Atom™ D525	x86_64	1.8 GHz	4 GiB	Asus AT5NM10T-I	Fedora 22	4.1.4	gcc-5.1.1	
<input type="radio"/> raspberry-pi-2	✓	Broadcom BCM2709	armv7l	900 MHz	1 GiB	Raspberry Pi 2 Model B	Raspbian Jessie	4.1.6	gcc-4.9.2	
<input type="radio"/> s-desktop	✓	Intel® Core™ i5-2400	x86_64	3.1 GHz	4 GiB	Asus P8Z68-V	Fedora 22	4.1.4	gcc-5.1.1	
<input type="radio"/> satellite-a205	✓	Intel® Celeron® Processor 540	x86_64	1.86 GHz	1 GiB	Toshiba Satellite A205-S5805	Fedora 21	4.1.6	gcc-4.9.2	

# COMPRESSION RATIO VS. DECOMPRESSION SPEED



Transfer Speed

PRESETS ▾

125

MIB/S ▾

Visible Items

Within 25% of no compression ▾

Sort



Total time



Name

Direction



Decompression




Compression



Both

DRAW




**Transfer**  
+  
**Decompression**

# Compression & Networking

# Background

---



- network links are getting faster and faster but
  - many clients still connected by fairly slow links (mobile?)
  - people wish to send more and more data
  - battery life is a king problem

how can we make this transparent to the user?



# Two standard techniques

---

- **caching**: *“avoid sending the same object again”*
  - Done on the basis of “atomic” objects
  - Thus only works if objects are unchanged
  - How about objects that are slightly changed?
- **compression**: *“remove redundancy in transmitted data”*
  - avoid repeated substrings in transmitted data
  - can be extended to history of past transmissions (overhead)
  - How if the sender has never seen data at receiver ?

# Types of Techniques

---

- Common knowledge between sender & receiver
  - Unstructured file: **delta compression**
- “Partial” knowledge between sender & receiver
  - Unstructured files: **file synchronization**
  - Record-based data: **set reconciliation**

# Formalization

---

- **Delta compression** [*diff, zdelta, REBL,...*]
  - Compress file  $f$  deploying **known** file  $f'$
  - Compress a group of files
  - Speed-up web access by sending differences between the requested page and the ones available in cache
- **File synchronization** [*rsynch, zsync*]
  - Client updates **its** old file  $f_{old}$  with **new** file  $f_{new}$  **available** on a server
  - Mirroring, Shared Crawling, Content Distribution Network

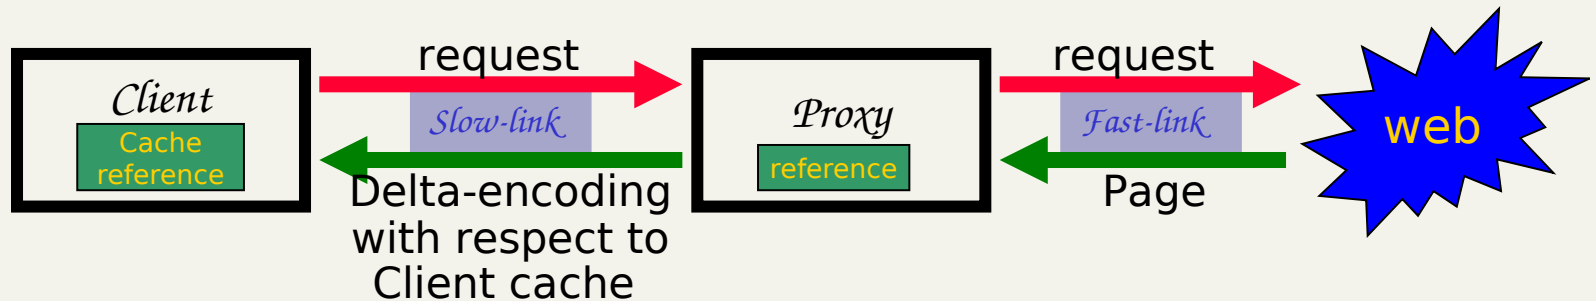
# Z-delta compression (one-to-one)

Problem: We have two files  $f_{\text{known}}$  (known to both parties) and  $f_{\text{new}}$  (is known only to the sender) and the goal is to compute a file  $f_d$  of minimum size such that  $f_{\text{new}}$  can be derived by the receiver from  $f_{\text{known}}$  and  $f_d$

- Assume that block **moves** and **copies** are allowed
- **Find an optimal covering set of  $f_{\text{new}}$  based on  $f_{\text{known}}$**
- **LZ77-scheme** provides an efficient, optimal solution
  - $f_{\text{known}}$  is “previously encoded text”, compress  $f_{\text{known}}f_{\text{new}}$  starting from  $f_{\text{new}}$
- **zdelta** is one of the best implementations
- Uses e.g. in Version control, Backups, and Transmission.

# Efficient Web Access

Dual proxy architecture: pair of proxies (client cache + proxy) located on each side of the slow link use a proprietary protocol to increase performance




- Use *zdelta* to reduce traffic:
  - Old version is available at both proxies (one on client cache, and one on proxy)
  - Restricted to pages already visited (30% cache hits), or UR

Small cache

# Cluster-based delta compression

Problem: We wish to compress a group of files  $F$

- Useful on a dynamic collection of web pages, back-ups, ...
- Apply *pairwise **zdelta***: find a *good reference* for each  $f \in F$
- Reduction to the Min Branching problem on DAGs
  - Build a (**complete?**) weighted graph  $G_F$ , nodes=files, weights= **zdelta**-size
  - Insert a dummy node connected to all, and weights are **gzip**-coding
  - Compute the directed spanning tree of min tot cost, covering  $G$ 's nodes.



		space	time
uncompr		30Mb	---
tgz		20%	linear
<i>THIS</i>		8%	quadratic

# Improvement (group of files)

Problem: Constructing  $G$  is very costly,  $n^2$  edge calculations (*zdelta exec*)

- We wish to exploit some pruning approach
  - **Collection analysis:** Cluster the files that appear similar and thus good candidates for zdelta-compression. Build a sparse weighted graph  $G'_F$  containing only edges between pairs of files in the same cluster
  - **Assign weights:** Estimate appropriate edge weights for  $G'_F$  thus saving zdelta execution. Nonetheless, strict  $n^2$  time

		space	time
uncompr		260Mb	---
tgz		12%	2 mins
<i>THIS</i>		8%	16 mins

# File Synchronization




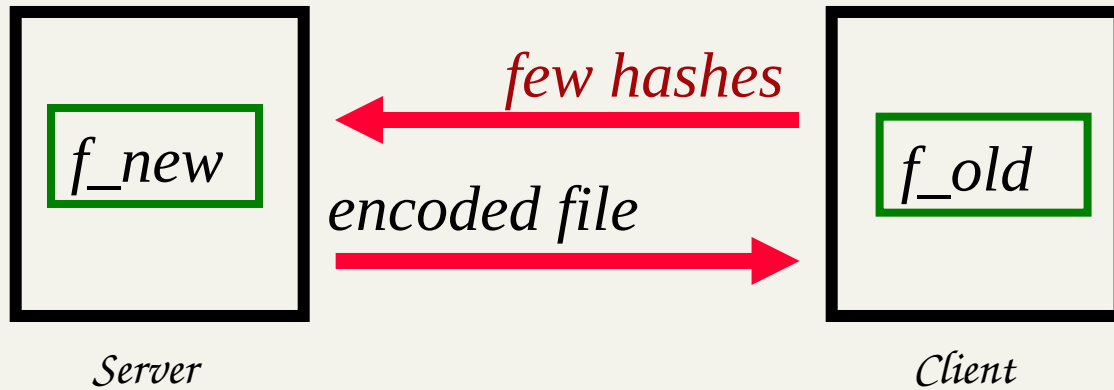
# File synch: The problem




- **client**
  - *request* to update an old file
  - Sends a sketch of the old one
- **server**
  - has new file but **does not know** the old file
  - Sends an update of  $f_{old}$  given its received sketch and  $f_{new}$
- *rsync*: file synch tool, distributed with Linux

Delta compression is a sort of local synch  
Since the server knows both files

# The rsync algorithm

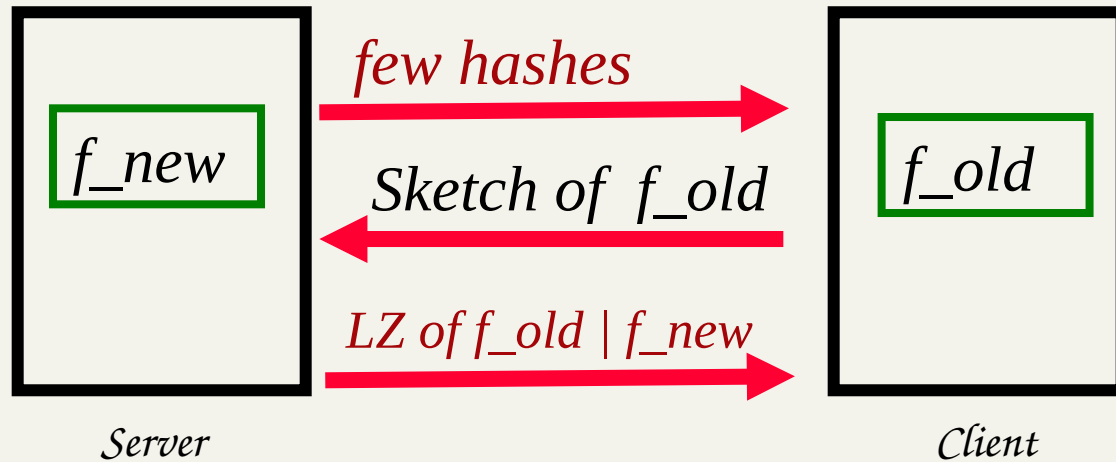


# The rsync algorithm (contd)



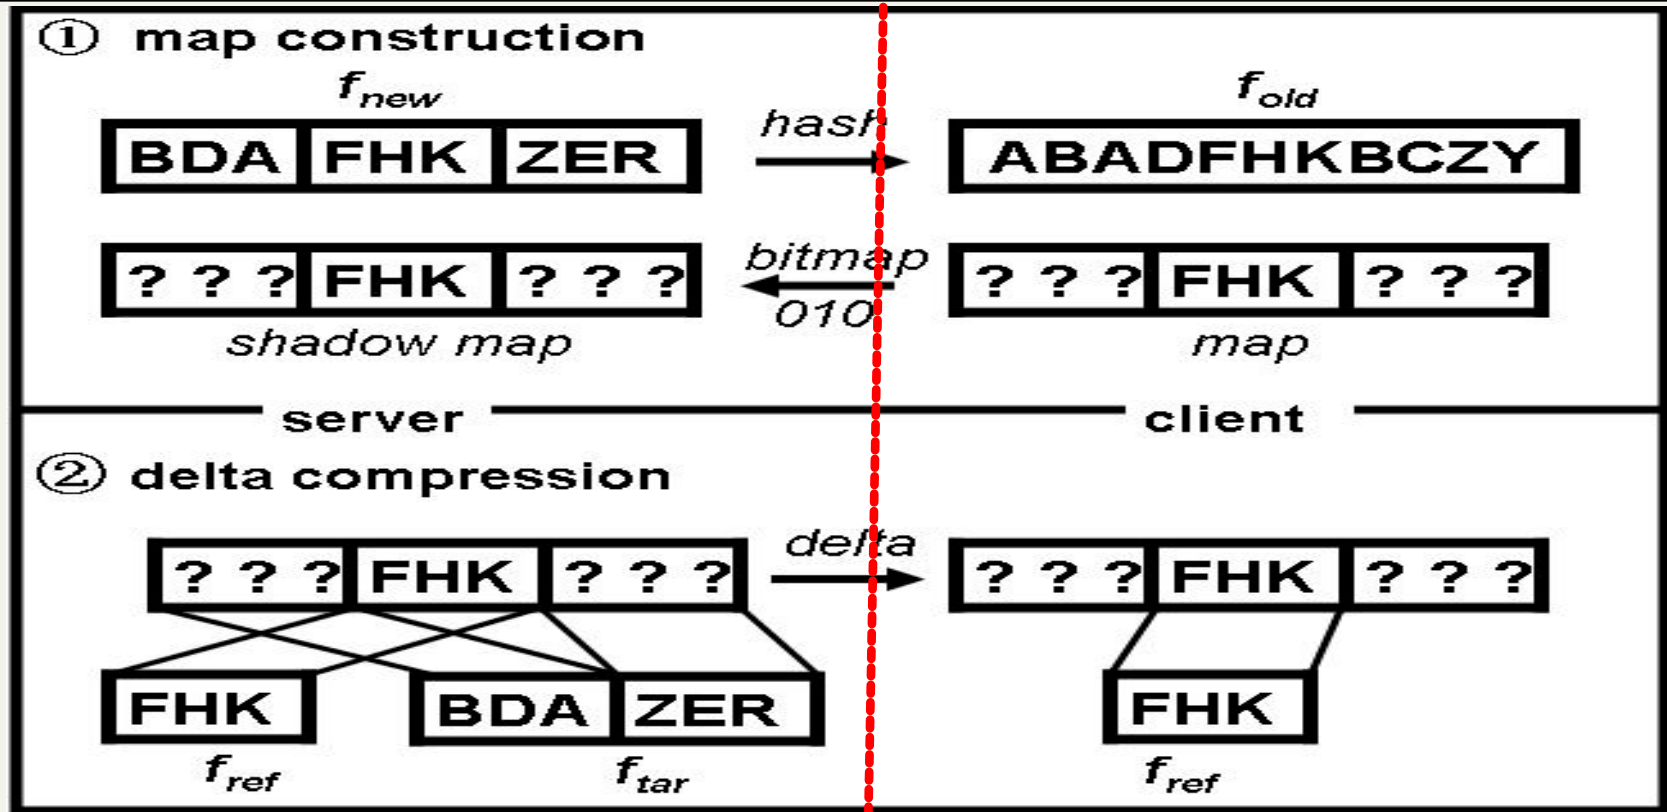
- simple, widely used, **single** roundtrip
- optimizations: 4-byte **rolling hash** + 2-byte **MD5**, **gzip** for literals
- choice of block size problematic (*default*:  $\max\{700, \sqrt{n}\}$  bytes)
- not good in theory: **granularity of changes may disrupt use of blocks**
- There is a **high load** on the server

# Minimize server load: **zsync**



- Server starts
- Client is assumed to be “good at CPU”: it makes the scan
- **Three** communication steps
- Hashes of  $f\_new$  precalculated, and stored in **.zsync**
- **Z-delta** is used to send  $f\_new$ , given a sketch of  $f\_old$

# Minimize server load: **zsync**



The hashes of the blocks can be precalculated and stored in .zsync file  
**Server** should be not overloaded. So it sends .zsync, clients checks them  
It is better suited to distribute multiple-files through network, given one .zsync

Bitmap = 3 bits because of #blocks in  $f_{new}$  are three